

Time_Series_Nifty_Bajaj_Deploy

Reading the market data of BAJAJFINSV stock and preparing a training dataset and validation dataset.

In [1]:

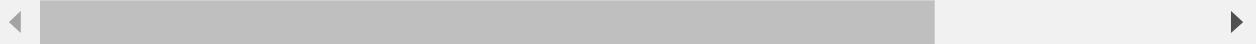
```
import pandas as pd
import numpy as np
```

In [2]:

```
df=pd.read_csv('BAJFINANCE.csv')
df.head()
```

Out[2]:

	Date	Symbol	Series	Prev Close	Open	High	Low	Last	Close	VWAP	Volume	Turnover
0	2000-01-03	BAJAUTOFIN	EQ	46.95	49.45	50.75	46.5	50.75	50.75	50.05	7600	3.803800e+10
1	2000-01-04	BAJAUTOFIN	EQ	50.75	53.20	53.20	47.9	48.00	48.10	48.56	5000	2.428000e+10
2	2000-01-05	BAJAUTOFIN	EQ	48.10	46.55	47.40	44.6	44.60	44.60	45.47	3500	1.591450e+10
3	2000-01-06	BAJAUTOFIN	EQ	44.60	43.50	46.00	42.1	46.00	45.25	44.43	6200	2.754750e+10
4	2000-01-07	BAJAUTOFIN	EQ	45.25	48.00	48.00	42.0	42.90	42.90	44.44	3500	1.555550e+10



In [3]:

```
df.set_index('Date', inplace=True)
```

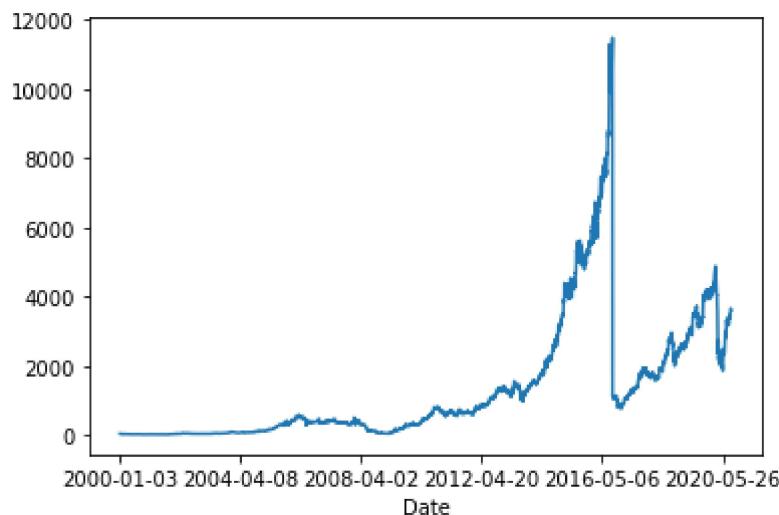
In []:

Plotting the target variable VWAP over time

In [4]:

```
df['VWAP'].plot()
```

Out[4]: <AxesSubplot:xlabel='Date'>



so u can observe here some kind of Seasonality

Feature Engineering Almost every time series problem will have some external features or some internal feature engineering to help the model. Let's add some basic features like lag values of available numeric features that are widely used for time series problems. Since we need to predict the price of the stock for a day, we cannot use the feature values of the same day since they will be unavailable at actual inference time. We need to use statistics like mean, standard deviation of their lagged values. We will use three sets of lagged values, one previous day, one looking back 7 days and another looking back 30 days as a proxy for last week and last month metrics.

Data Pre-Processing

```
In [5]: df.shape
```

```
Out[5]: (5070, 14)
```

```
In [6]: df.isna().sum()
```

```
Out[6]: Symbol          0
Series           0
Prev Close       0
Open             0
High             0
Low              0
Last             0
Close            0
VWAP             0
Volume           0
Turnover         0
Trades          2779
Deliverable Volume 446
%Deliverable     446
dtype: int64
```

```
In [8]: df.dropna(inplace=True)
```

```
In [9]: df.isna().sum()
```

```
Out[9]: Symbol          0
```

```
Series          0
Prev Close     0
Open           0
High           0
Low            0
Last           0
Close          0
VWAP           0
Volume         0
Turnover       0
Trades          0
Deliverable Volume 0
%DDeliverble   0
dtype: int64
```

In [10]: df.shape

Out[10]: (2291, 14)

In [11]: data=df.copy()

In [12]: data.dtypes

```
Out[12]: Symbol          object
Series           object
Prev Close      float64
Open             float64
High            float64
Low              float64
Last             float64
Close            float64
VWAP            float64
Volume          int64
Turnover        float64
Trades           float64
Deliverable Volume float64
%DDeliverble    float64
dtype: object
```

In [13]: data.columns

```
Out[13]: Index(['Symbol', 'Series', 'Prev Close', 'Open', 'High', 'Low', 'Last',
 'Close', 'VWAP', 'Volume', 'Turnover', 'Trades', 'Deliverable Volume',
 '%Deliverble'],
 dtype='object')
```

In [14]: lag_features=['High','Low','Volume','Turnover','Trades']
window1=3
window2=7

In [15]: for feature in lag_features:
 data[feature+'rolling_mean_3']=data[feature].rolling(window=window1).mean()
 data[feature+'rolling_mean_7']=data[feature].rolling(window=window2).mean()

In [16]: for feature in lag_features:

```
data[feature+'rolling_std_3']=data[feature].rolling(window=window1).std()
data[feature+'rolling_std_7']=data[feature].rolling(window=window2).std()
```

In [17]: `data.head()`

Out[17]:

	Symbol	Series	Prev Close	Open	High	Low	Last	Close	VWAP	Volume	...	Highrolli
	Date											
2011-06-01	BAJFINANCE	EQ	616.70	617.00	636.50	616.00	627.00	631.85	627.01	6894	...	
2011-06-02	BAJFINANCE	EQ	631.85	625.00	638.90	620.00	634.00	633.45	636.04	2769	...	
2011-06-03	BAJFINANCE	EQ	633.45	625.15	637.80	620.00	623.00	625.00	625.09	51427	...	
2011-06-06	BAJFINANCE	EQ	625.00	620.00	641.00	611.35	611.35	614.00	616.03	5446	...	
2011-06-07	BAJFINANCE	EQ	614.00	604.00	623.95	604.00	619.90	619.15	617.73	5991	...	

5 rows × 34 columns



In [18]: `data.columns`

Out[18]: `Index(['Symbol', 'Series', 'Prev Close', 'Open', 'High', 'Low', 'Last', 'Close', 'VWAP', 'Volume', 'Turnover', 'Trades', 'Deliverable Volume', '%Deliverable', 'Highrolling_mean_3', 'Highrolling_mean_7', 'Lowrolling_mean_3', 'Lowrolling_mean_7', 'Volumerolling_mean_3', 'Volumerolling_mean_7', 'Turnoverrolling_mean_3', 'Turnoverrolling_mean_7', 'Tradesrolling_mean_3', 'Tradesrolling_mean_7', 'Highrolling_std_3', 'Highrolling_std_7', 'Lowrolling_std_3', 'Lowrolling_std_7', 'Volumerolling_std_3', 'Volumerolling_std_7', 'Turnoverrolling_std_3', 'Turnoverrolling_std_7', 'Tradesrolling_std_3', 'Tradesrolling_std_7'], dtype='object')`

In [19]: `data.shape`

Out[19]: `(2291, 34)`

In [20]: `data.isna().sum()`

Out[20]:

Symbol	0
Series	0
Prev Close	0
Open	0
High	0
Low	0
Last	0
Close	0

```
VWAP          0
Volume        0
Turnover      0
Trades         0
Deliverable Volume 0
%Deliverble   0
Highrolling_mean_3 2
Highrolling_mean_7 6
Lowrolling_mean_3 2
Lowrolling_mean_7 6
Volumerolling_mean_3 2
Volumerolling_mean_7 6
Turnoverrolling_mean_3 2
Turnoverrolling_mean_7 6
Tradesrolling_mean_3 2
Tradesrolling_mean_7 6
Highrolling_std_3 2
Highrolling_std_7 6
Lowrolling_std_3 2
Lowrolling_std_7 6
Volumerolling_std_3 2
Volumerolling_std_7 6
Turnoverrolling_std_3 2
Turnoverrolling_std_7 6
Tradesrolling_std_3 2
Tradesrolling_std_7 6
dtype: int64
```

In [21]: `data.dropna(inplace=True)`

In [22]: `data.columns`

Out[22]: `Index(['Symbol', 'Series', 'Prev Close', 'Open', 'High', 'Low', 'Last',
'Close', 'VWAP', 'Volume', 'Turnover', 'Trades', 'Deliverable Volume',
'%Deliverble', 'Highrolling_mean_3', 'Highrolling_mean_7',
'Lowrolling_mean_3', 'Lowrolling_mean_7', 'Volumerolling_mean_3',
'Volumerolling_mean_7', 'Turnoverrolling_mean_3',
'Turnoverrolling_mean_7', 'Tradesrolling_mean_3',
'Tradesrolling_mean_7', 'Highrolling_std_3', 'Highrolling_std_7',
'Lowrolling_std_3', 'Lowrolling_std_7', 'Volumerolling_std_3',
'Volumerolling_std_7', 'Turnoverrolling_std_3', 'Turnoverrolling_std_7',
'Tradesrolling_std_3', 'Tradesrolling_std_7'],
dtype='object')`

In [23]: `ind_features=['Highrolling_mean_3', 'Highrolling_mean_7',
'Lowrolling_mean_3', 'Lowrolling_mean_7', 'Volumerolling_mean_3',
'Volumerolling_mean_7', 'Turnoverrolling_mean_3',
'Turnoverrolling_mean_7', 'Tradesrolling_mean_3',
'Tradesrolling_mean_7', 'Highrolling_std_3', 'Highrolling_std_7',
'Lowrolling_std_3', 'Lowrolling_std_7', 'Volumerolling_std_3',
'Volumerolling_std_7', 'Turnoverrolling_std_3', 'Turnoverrolling_std_7',
'Tradesrolling_std_3', 'Tradesrolling_std_7']`

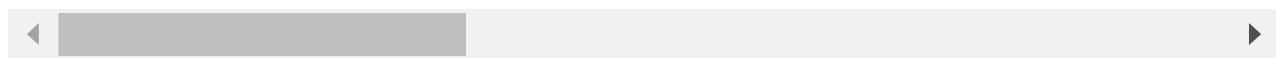
In [24]: `training_data=data[0:1800]
test_data=data[1800:]`

In [25]: `training_data`

Out[25]:

	Symbol	Series	Prev Close	Open	High	Low	Last	Close	VWAP	Volume	...
Date											
2011-06-09	BAJFINANCE	EQ	635.60	639.80	647.00	630.00	630.00	631.10	638.27	31252	...
2011-06-10	BAJFINANCE	EQ	631.10	641.85	648.25	618.55	621.10	622.20	634.16	30885	...
2011-06-13	BAJFINANCE	EQ	622.20	616.00	627.85	616.00	622.75	624.95	622.92	3981	...
2011-06-14	BAJFINANCE	EQ	624.95	625.00	628.95	619.95	621.20	622.10	625.35	5597	...
2011-06-15	BAJFINANCE	EQ	622.10	612.00	623.00	598.10	605.00	601.70	606.90	12590	...
...
2018-09-04	BAJFINANCE	EQ	2724.05	2724.00	2777.65	2683.50	2748.00	2746.30	2726.23	2606992	...
2018-09-05	BAJFINANCE	EQ	2746.30	2740.15	2764.80	2668.00	2704.45	2716.90	2712.53	1728455	...
2018-09-06	BAJFINANCE	EQ	2716.90	2729.00	2731.50	2671.40	2672.20	2684.10	2695.89	1147879	...
2018-09-07	BAJFINANCE	EQ	2684.10	2698.40	2751.40	2672.60	2745.00	2744.20	2716.32	1264436	...
2018-09-10	BAJFINANCE	EQ	2744.20	2732.00	2738.00	2596.00	2607.60	2615.65	2655.39	1570179	...

1800 rows × 34 columns



In []:

In [26]:

!pip install pmdarima

```
Collecting pmdarima
  Downloading pmdarima-1.8.5-cp38-cp38-win_amd64.whl (602 kB)
Requirement already satisfied: numpy>=1.19.3 in c:\anaconda\lib\site-packages (from pmdarima) (1.20.1)
Requirement already satisfied: joblib>=0.11 in c:\anaconda\lib\site-packages (from pmdarima) (1.0.1)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in c:\anaconda\lib\site-packages (from pmdarima) (52.0.0.post20210125)
Requirement already satisfied: scipy>=1.3.2 in c:\anaconda\lib\site-packages (from pmdarima) (1.6.2)
Requirement already satisfied: scikit-learn>=0.22 in c:\anaconda\lib\site-packages (from pmdarima) (0.24.1)
Requirement already satisfied: Cython!=0.29.18,>=0.29 in c:\anaconda\lib\site-packages (from pmdarima) (0.29.23)
Requirement already satisfied: urllib3 in c:\anaconda\lib\site-packages (from pmdarima) (1.26.4)
```

```
Requirement already satisfied: pandas>=0.19 in c:\anaconda\lib\site-packages (from pmdarima) (1.2.4)
Requirement already satisfied: statsmodels!=0.12.0,>=0.11 in c:\anaconda\lib\site-packages (from pmdarima) (0.12.2)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\anaconda\lib\site-packages (from pandas>=0.19->pmdarima) (2.8.1)
Requirement already satisfied: pytz>=2017.3 in c:\anaconda\lib\site-packages (from pandas>=0.19->pmdarima) (2021.1)
Requirement already satisfied: six>=1.5 in c:\anaconda\lib\site-packages (from python-dateutil>=2.7.3->pandas>=0.19->pmdarima) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\anaconda\lib\site-packages (from scikit-learn>=0.22->pmdarima) (2.1.0)
Requirement already satisfied: patsy>=0.5 in c:\anaconda\lib\site-packages (from statsmodels!=0.12.0,>=0.11->pmdarima) (0.5.1)
Installing collected packages: pmdarima
Successfully installed pmdarima-1.8.5
```

In [27]: `from pmdarima import auto_arima`

In [28]: `import warnings
warnings.filterwarnings('ignore')`

In [29]: `model=auto_arima(y=training_data['VWAP'], exogenous=training_data[ind_features], trace=True)`

```
Performing stepwise search to minimize aic
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=20931.563, Time=3.41 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=20925.246, Time=2.45 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=20926.371, Time=3.41 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=20926.343, Time=2.96 sec
ARIMA(0,0,0)(0,0,0)[0] : AIC=32616.914, Time=2.44 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=20929.256, Time=4.29 sec
```

```
Best model: ARIMA(0,0,0)(0,0,0)[0] intercept
Total fit time: 19.104 seconds
```

In [30]: `model.fit(training_data['VWAP'], training_data[ind_features])`

Out[30]: `ARIMA(order=(0, 0, 0), scoring_args={}, suppress_warnings=True)`

In [31]: `forecast=model.predict(n_periods=len(test_data), exogenous=test_data[ind_features])`

In [32]: `test_data['Forecast_ARIMA']=forecast`

In [34]: `test_data[['VWAP','Forecast_ARIMA']].plot(figsize=(14,7))`

Out[34]: <AxesSubplot:xlabel='Date'>



The Auto ARIMA model seems to do a fairly good job in predicting the stock price

In []:

Checking Accuracy of our model

In [35]:

```
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

In [36]:

```
np.sqrt(mean_squared_error(test_data['VWAP'], test_data['Forecast_ARIMA']))
```

Out[36]:

```
187.91287668231973
```

In [37]:

```
mean_absolute_error(test_data['VWAP'], test_data['Forecast_ARIMA'])
```

Out[37]:

```
124.74418035631113
```

In []:

In []: